

# SYSTEMS PROGRAMMING AND COMPUTER ARCHITECTURE

## Assignment 12: Cache Coherence & Page Table Programming

Assigned on: 12th December 2018  
Due by: 18th December 2018, 23:59

### Part I: Pen & Paper Exercises

#### Question 1

Consider the system in the figure below, where memory write operations are performed in the write-back mode.




Figure 1: Model of a multi-core architecture with two processors  $P_A$  and  $P_B$ , their respective caches  $C_A$  and  $C_B$ , connected to main memory  $M$  over a system bus.

At first, all of the cache-lines are marked "invalid". The value 10 is stored in memory at the address  $X$ . The two processors execute two sequences of instructions described in the following table.

$n$	$P_A$	$P_B$	Comment
1	mov $(X), r1$		$P_A: r1 := (X)$
2		mov $(X), r4$	$P_B: r4 := (X)$
3	mov $\$0, (X)$		$P_A: (X) := 0$
4		mov $(X), r5$	$P_B: r5 := (X)$
5		mov $\$20, (X)$	$P_B: (X) := 20$

Table 1: Operations executed by processor  $P_A$  and  $P_B$ .

- a) In the table below, describe the sequence of operations that ensure cache coherence according to the MSI protocol. In each step, assume the instruction has already been executed. For each phase specify the state of the cache lines, the value of X they are holding, and the value of X in main memory.

$n$	State $C_A$	Value $C_A$	State $C_B$	Value $C_B$	Value $M$
1					
2					
3					
4					
5					

- b) Now describe the same sequence but assume the architecture uses the MESI protocol.

$n$	State $C_A$	Value $C_A$	State $C_B$	Value $C_B$	Value $M$
1					
2					
3					
4					
5					

## Question 2

Consider three cores equipped with a direct mapped, write-back cache of 128 bytes capacity using a cache-line size of 8 bytes. These machines executes "load-store" sequences in the order described in the table below. The size of loads and stores is two bytes.  $n$  is the time when instructions are executed. X contains address  $0xA0C0$ .

$n$	$P_1$	$P_2$	$P_3$
1	ld r1, [X]	...	...
2	add r1, 1	...	...
3	st r1, [X]	...	...
4	...	ld r2, [X+2]	...
5	...	and r2, 0FH	...
6	...	st r2, [X+2]	...
7	...	...	ld r3, [X+6]
8	...	...	sub r3, r1, r5
9	...	...	st r3, [X+6]

- Which cache-line will be used by the three processors?
- Describe MSI protocol transition for the caches of each processor in every step, after the instruction has been executed. Assume all cache-lines are marked as invalid in the beginning.

$n$	State $C_1$	State $C_2$	State $C_3$
1			
2			
3			
4			
5			
6			
7			
8			
9			

## Hand In Instructions for Part I

This part is a paper exercise. If you want your solution to be revised please hand it in during your exercise class on the due date.

## Part II: Programming

In this assignment you will implement basic virtual memory functionality in C based on the structure of a x86\_64 page-table. x86\_64 has a four-level page table and these are referred to as the page map level 4, page directory pointer table, page directory and page table respectively.

The following figure shows the bit layout for the different levels of paging structures. (In our case we consider `MAXPHYADDR`, the width of the physical addresses, to be 48 bits.)

6	6	6	6	5	5	5	5	5	5	5	5	5		M	M-1		3	3	3	3	2	2	2	2	2	2	2	2	1	1	1	1	1	1	1	1	0		
3	2	1	0	9	8	7	6	5	4	3	2	1					2	1	0	9	8	7	6	5	4	3	2	1	0	9	8	7	6	5	4	3	2	1	0
Reserved														Address of PML4 table														Ignored				P C W D T	P I g n.	CR3					
X D	Ignored												Rsvd.	Address of page-directory-pointer table														Ign.	R s v d	I g n	A	P C W D T	P U S	R / W	1	PML4E: present			
Ignored																												0	PML4E: not present										
X D	Ignored												Rsvd.	Address of 1GB page frame				Reserved								P A T	Ign.	G 1	D A	P C W D T	P U S	R / W	1	PDPTE: 1GB page					
X D	Ignored												Rsvd.	Address of page directory														Ign.	0	I g n	A	P C W D T	P U S	R / W	1	PDPTE: page directory			
Ignored																												0	PDPTE: not present										
X D	Ignored												Rsvd.	Address of 2MB page frame								Reserved				P A T	Ign.	G 1	D A	P C W D T	P U S	R / W	1	PDE: 2MB page					
X D	Ignored												Rsvd.	Address of page table														Ign.	0	I g n	A	P C W D T	P U S	R / W	1	PDE: page table			
Ignored																												0	PDE: not present										
X D	Ignored												Rsvd.	Address of 4KB page frame														Ign.	G A T	D A	P C W D T	P U S	R / W	1	PTE: 4KB page				
Ignored																												0	PTE: not present										


Figure 2: Formats of Paging-Structure Entries<sup>1</sup>

The architecture manuals contain a thorough description of the flags used in figure above and the most pertinent fields are summarized below:

Bit Position	Contents
(M-1):12	Physical address of the page table or page referenced by this entry
7 (PS)	Page Size; must be one if this refers to a large/huge page, otherwise this entry references a page table or regular page
1 (R/W)	Read/Write; if 0, writes may not be allowed
0 (P)	Present; marks whether the entry is present or not

<sup>1</sup>This figure is taken from Figure 4-11, Page 4-28 of the Intel® 64 and IA-32 Architectures Software Developer's Manual, Volume 3A: System Programming Guide, Part 1

The structure of a 48-bit virtual address is:



VPN 1 – 4 are the virtual page numbers used as lookup indices for the four levels of page table, whereas VPO is the virtual page offset which addresses the byte within the page.

## Task 1

We prepared a skeleton that can be used as a basis for the implementation. You will find it on the course website as a tar archive. Download and extract that file.

As a first step, read through understand the type definitions which can be found in `<paging.h>`:

Type	Base type	Description
<code>vaddr_t</code>	<code>uint64_t</code>	Virtual address
<code>paddr_t</code>	<code>uint64_t</code>	Physical address
<code>pml4e_t</code>	<code>union</code>	Entry in the page map level 4
<code>ppde_t</code>	<code>union</code>	Entry in the page directory pointer table (level 3)
<code>pde_t</code>	<code>union</code>	Entry in the page directory (level 2)
<code>pte_t</code>	<code>union</code>	Entry in the page table (level 1)

These types define the bit layouts of page directory and page table entries. The terminology we use here is that a “page directory entry” is referring to another page table or page directory while a page table entry is referring to a mapped region of physical memory.

Note that we are also using C bitfields here. Bitfields are structs where we specify the bit widths of different fields using `<bits>` after the field name. This way you can access some bits without having to manually shift and mask. Additionally we’re wrapping the bitfield and a `uint64_t` in an union to be able to easily convert between raw values and structured values of page table and page directory entries.

Next, implement functions `parse_virt_addr`, `get_pml4e`, `get_pdpe`, `get_pde`, `get_pte`. The first of these should extract the page table indices from a virtual address; the remainder are utility functions to read a page table entry and walk to the next level of the page table.

## Task 2

Have a look at the skeletons for functions `map`, `map_large` and `map_huge`. These functions map physical pages into the virtual address space by inserting entries into the respective levels of the page tables. They also set some attributes in these entries. At this point, you only have to handle the read/write permissions (i.e. if at least one page is writable, the entire path leading to it should also be marked writable).

For regular 4 KiB pages, `map` needs to touch all four levels of the page-table hierarchy, whereas `map_large` (for 2 MiB mappings) and `map_huge` (for 1 GiB mappings) only modify the content of the second- and third-level page tables (page directory and page directory pointer table respectively). Mapping a large or huge page works by pointing the respective page table entry directly to a 2 MiB or 1 GiB chunk of memory. Note that the memory itself needs to be appropriately aligned, as only the uppermost 18/27/36-bits of the physical address can be stored in the different levels of the page table entries. The PS bit in the second and third-level entries is used to distinguish huge, large and regular pages. Implement these three functions.

In case the page-directory or the appropriate page-table does not yet exist, you will have to allocate some space for it. Have a look at and implement the `alloc_table` function. Remember: Both the

beginning of the page-table and the page-directory data-structures need to be page-aligned. Do not forget to implement *free\_pagetable* to free all allocated memory again.

### Task 3

Next, you are going to implement the *unmap* operation. This routine takes a virtual address and removes the respective entries by setting the value to zero (including the present bit). It should work for both, 4 KiB, 2 MiB and 1 GiB pages.

### Evaluation

You can use the *./correctness* script in the handout archive to verify your solution. The program will run your page table implementation and compare it to a reference solution. *dropAddresses* is executed on your output first to remove all dynamic addresses from the page table entries. (These are the addresses stored in the page directory for the pages themselves.) Then, the Unix *diff* command is executed to compare the files.

### Hints

Run *make* after extracting the archive we provide on the website to compile your page table manipulation program. gcc will link against a static library *libdump.a* required to dump your page table in a format we can parse with the *./correctness* script. Your program will be compiled in 64-bit mode. As always, you should install the *build-essential* package which contains the compiler and build tools if you have your own environment.

### Hand In Instructions for Part II

Upload your source files to a subfolder **assignment12** of your SVN folder.