

Systems Programming and Computer Architecture


Exercise Session 14

“The last one”

- Assignment 12
- Devices recap
- Assignment 13

Exercise 12


Question 1a: MSI



n	P_A	P_B	Comment
1	mov (X),r1		P_A : r1 := (X)
2		mov (X),r4	P_B : r4 := (X)
3	mov \$0,(X)		P_A : (X) := 0
4		mov (X),r5	P_B : r5 := (X)
5		mov \$20,(X)	P_B : (X) := 20

n	State Ca	Value Ca	State Cb	Value Cb	Value M
1	S	10	I	—	10
2	S	10	S	10	10
3	M	0	I	—	10
4	S	0	S	0	0
5	I	—	M	20	0

Question 1b: MESI



n	P_A	P_B	Comment
1	mov (X),r1		P_A : r1 := (X)
2		mov (X),r4	P_B : r4 := (X)
3	mov \$0,(X)		P_A : (X) := 0
4		mov (X),r5	P_B : r5 := (X)
5		mov \$20,(X)	P_B : (X) := 20

n	State Ca	Value Ca	State Cb	Value Cb	Value M
1	E	10	I	—	10
2	S	10	S	10	10
3	M	0	I	—	10
4	S	0	S	0	0
5	I	—	M	20	0

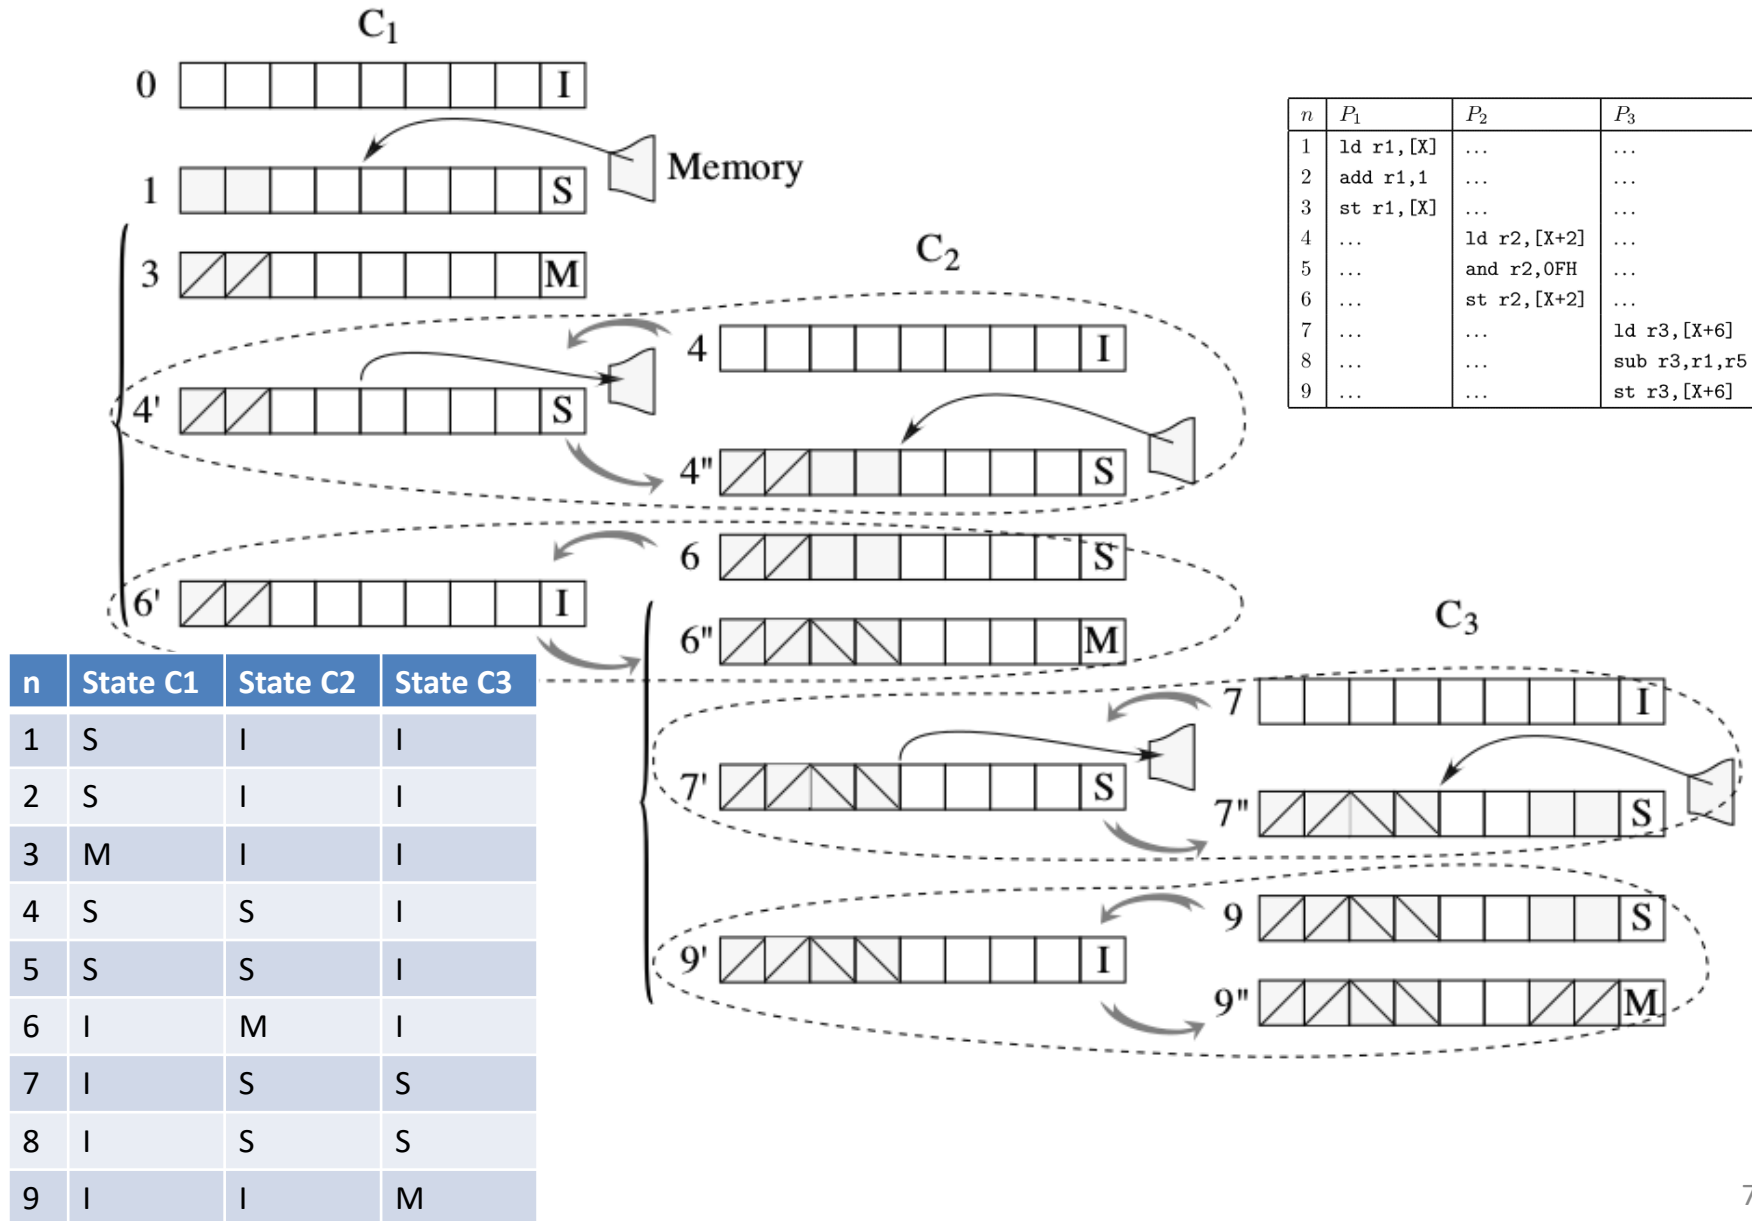

Only
difference
between
MESI and
MSI

Question 2

- 2a: 0xA0C0 maps to index 8 (bits 6-3).
- 3 cores
- 128b direct-mapped write-back cache, 8b cache line size
- $X = 0xA0C0$

n	P_1	P_2	P_3
1	ld r1, [X]
2	add r1, 1
3	st r1, [X]
4	...	ld r2, [X+2]	...
5	...	and r2, 0FH	...
6	...	st r2, [X+2]	...
7	ld r3, [X+6]
8	sub r3, r1, r5
9	st r3, [X+6]

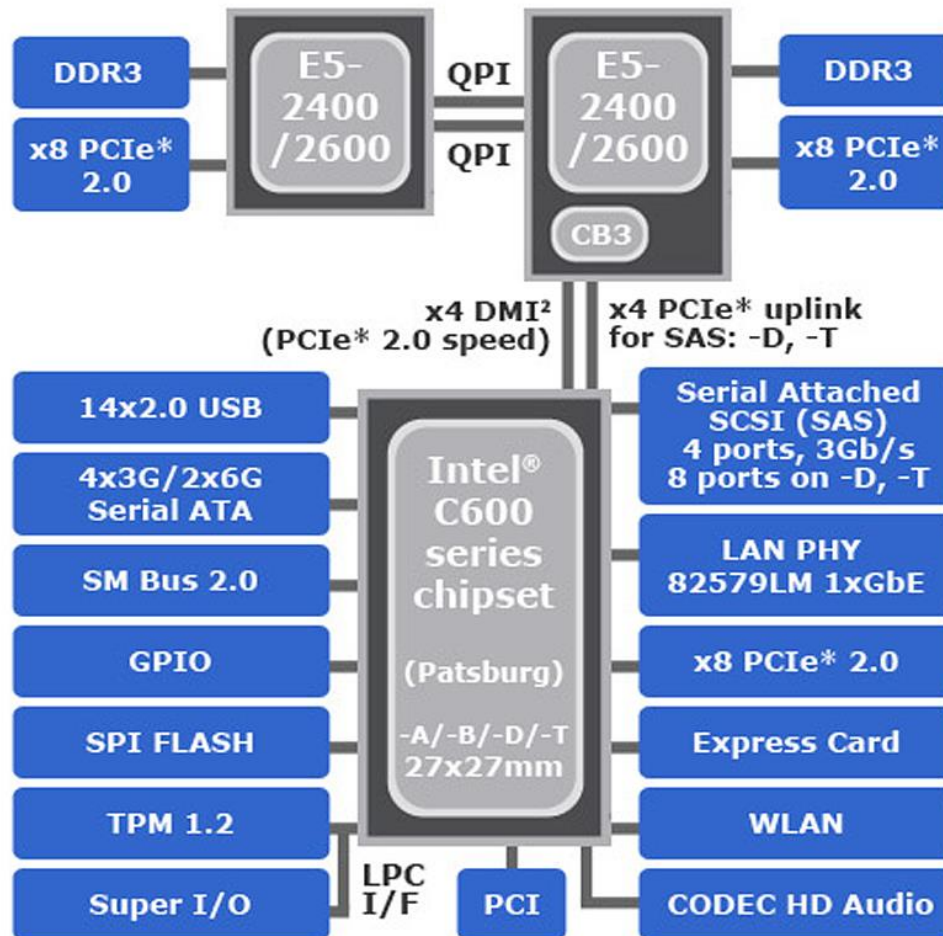
Question 2b: MSI



Devices Recap

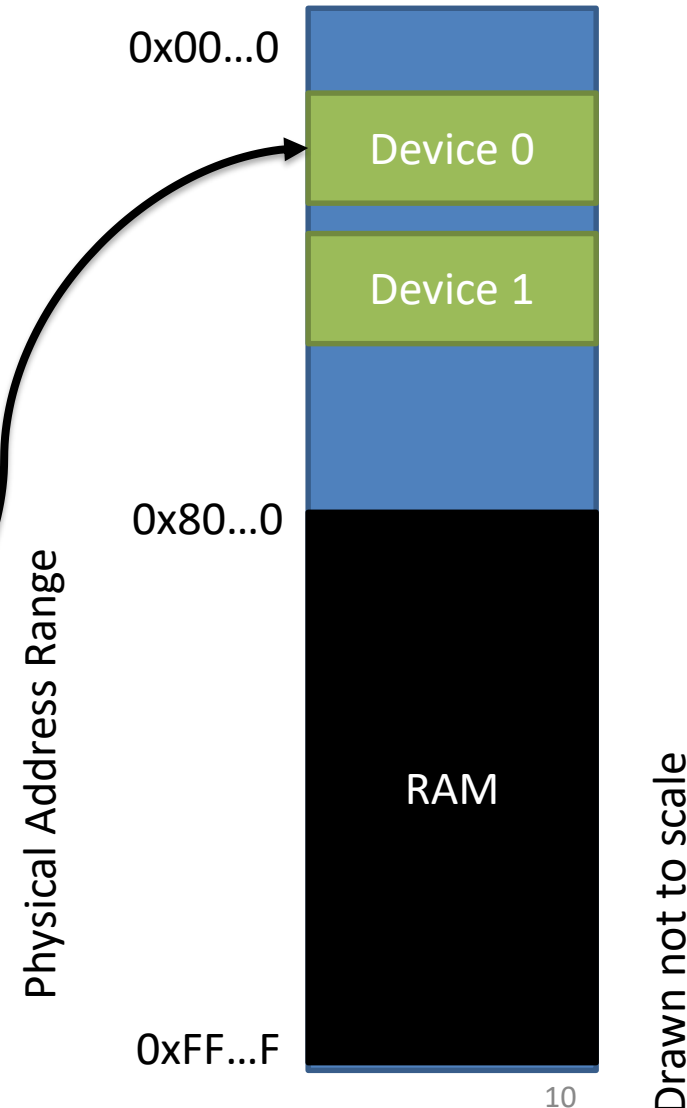



Devices



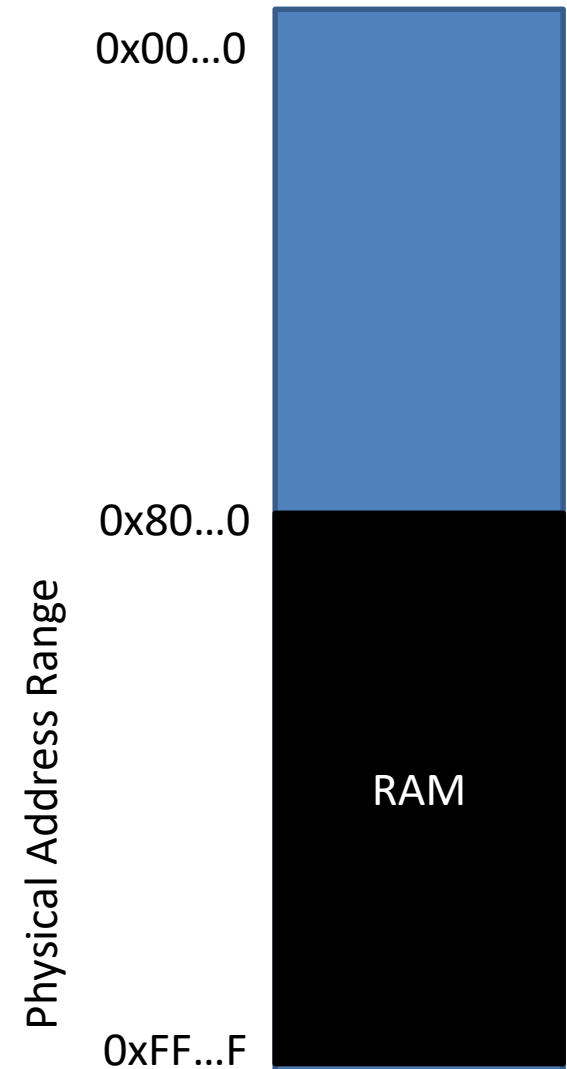

How to access devices?

- Memory mapped
 - Write & Read to devices as you have done with normal memory (movl...)
 - Use of the MMU & Page Tables



How to access devices?

- IO Ports
 - Special address space (16 bit)
 - Write to an IO port



Memory Mapped Devices

- Device represented as [Base Address, Length]
 - Base address refers to the physical address where the device starts
 - cf: your beginning of the stack frame
 - Length total occupied memory region by that device (actually used by registers may be lower)
 - cf: the size of your stack frame
 - Set of registers within Base, Base+Length
 - cf: your variables on the stack

Devices are **NOT** memory

- Contents of register may change unexpectedly
 - Data received
- Writing to a register trigger actions
 - Shutdown device / machine
 - Perform reset

Too simple UART driver

```
1. #define UART_BASE 0x12345000
2. #define UART_THR (UART_BASE + 0)
3. #define UART_RBR (UART_BASE + 4)
4. #define UART_LSR (UART_BASE + 8)

5. void serial_putc(char c)
6. {
7.     char *lsr = (char *) UART_LSR;
8.     char *thr = (char *) UART_THR;

9.     // Wait until FIFO can hold more chars
10.    while((*lsr & 0x20) == 0);
11.
12.    *thr = c
13. }
```

- What's the problem here?

Too simple UART driver

```
1. #define UART_BASE 0x12345000
2. #define UART_THR (UART_BASE + 0)
3. #define UART_RBR (UART_BASE + 4)
4. #define UART_LSR (UART_BASE + 8)

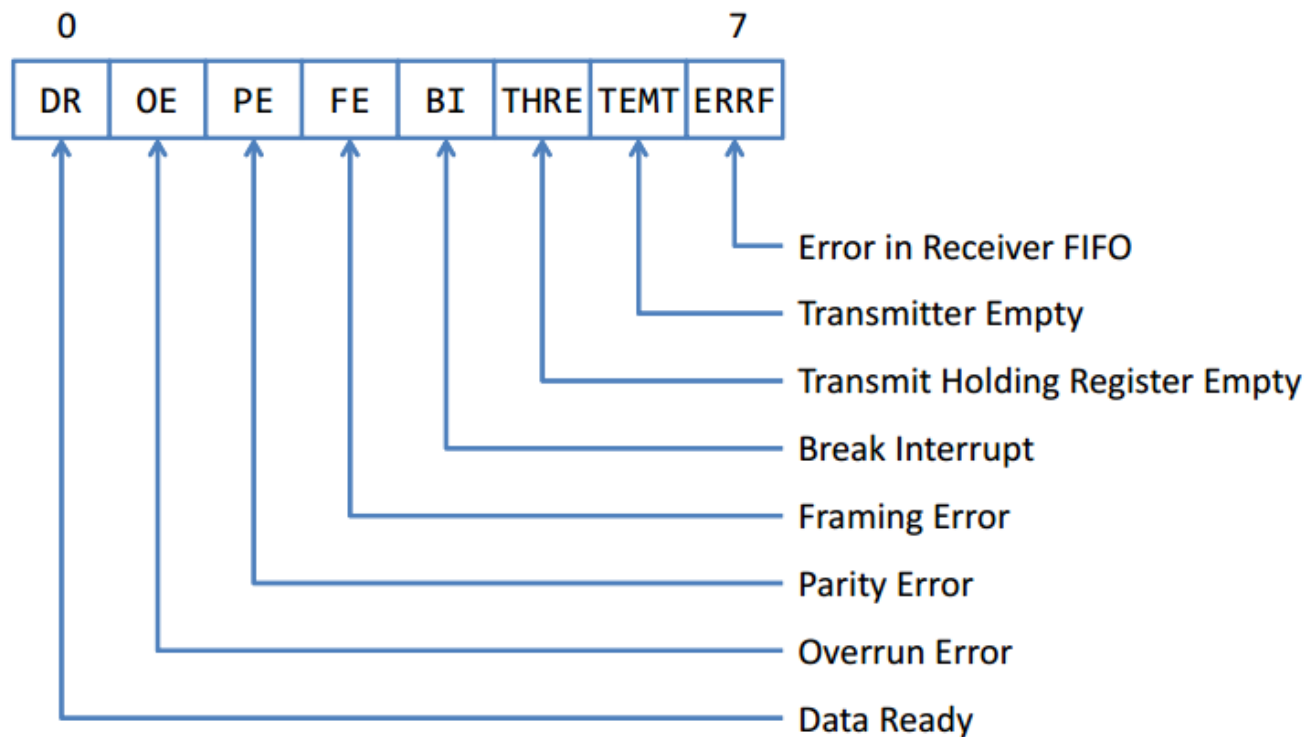
5. void serial_putc(char c)
6. {
7.     volatile char *lsr = (char *) UART_LSR;
8.     volatile char *thr = (char *) UART_THR;

9.     // Wait until FIFO can hold more chars
10.    while((*lsr & 0x20) == 0);
11.
12.    *thr = c
13. }
```

- Compiler does not know this is a device register!
- Loop gets optimized away!
- Add: volatile

Register Contents

- Each bit may have a different meaning
- Lots of configurations possible!



Devices

- Writing device drivers is tedious
 - Setting a single bit wrong and the device does not work
 - Debugging is hard: Likely to set a bit wrong due to a wrong shift & mask.



Devices & Caches

- Device registers cannot be cached due to inconsistency problem i.e. register content changes w/o CPU write!
- What about cache lines? Would overwrite other register values when writing back
- Set the “no-cache” flag in the page table entry

Interrupts

```
1. #define UART_BASE 0x12345000
2. #define UART_THR (UART_BASE + 0)
3. #define UART_RBR (UART_BASE + 4)
4. #define UART_LSR (UART_BASE + 8)

5. void serial_putc(char c)
6. {
7.     volatile char *lsr = (char *) UART_LSR;
8.     volatile char *thr = (char *) UART_THR;

9.     // Wait until FIFO can hold more chars
10.    while((*lsr & 0x20) == 0);
11.
12.    *thr = c
13. }
```

- Any “problems” with that one?

Interrupts

```
1. #define UART_BASE 0x12345000
2. #define UART_THR (UART_BASE + 0)
3. #define UART_RBR (UART_BASE + 4)
4. #define UART_LSR (UART_BASE + 8)



5. void serial_putc(char c)
6. {
7.     volatile char *lsr = (char *) UART_LSR;
8.     volatile char *thr = (char *) UART_THR;

9.     // Wait until FIFO can hold more chars
10.    while((*lsr & 0x20) == 0);
11.
12.    *thr = c
13. }
```

- Register callback for an interrupt
- Activate interrupt


Polls register
=
wastes CPU cycles

DMA



DMA

- DMA is like device registers! Changes the contents w/o CPU write.
- Cannot “no-cache”:
Bad performance
- Explicitly invalidate
cache!



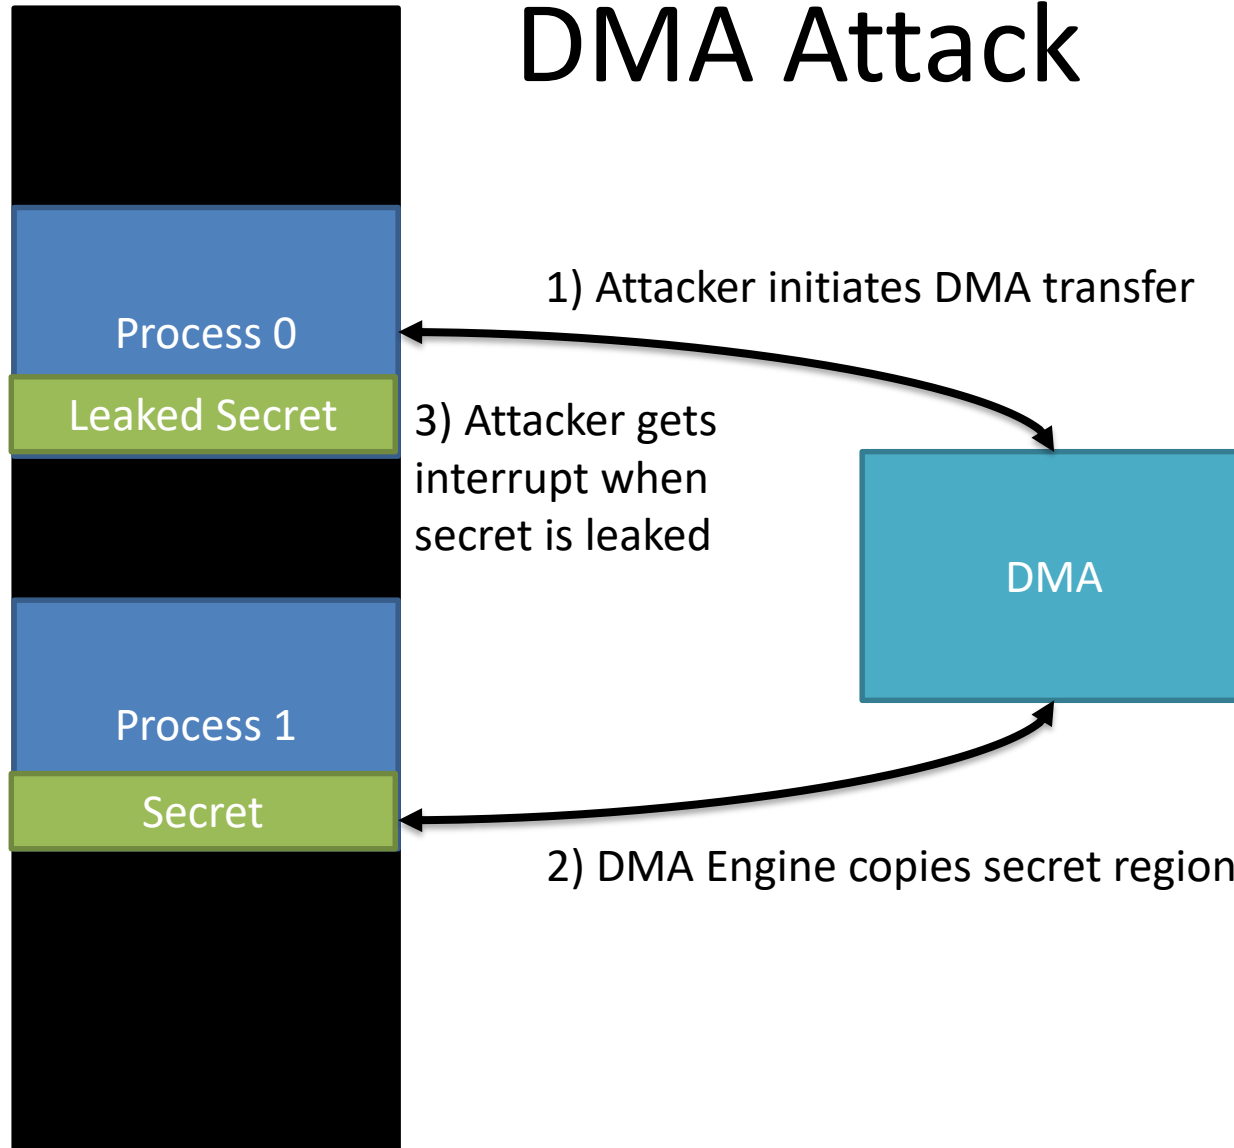
DMA

- Deal with physical addresses only!
 - Any problems with that?

DMA

- Deal with physical addresses only!
 - Any problems with that?
 - Programs deal with virtual addresses: need translation!
 - Physical Range may not be contiguous (gather/scatter)
- How about security?

DMA Attack



Attack done by external devices i.e. firewire / thunderbolt

http://en.wikipedia.org/wiki/DMA_attack

DMA & Cache Problems

- On DMA read:
 - Before: Flush the cache to update main memory
 - After: Invalidate the cache to avoid old values seen
- On DMA write:
 - Before: flush or invalidate the cache to update main memory
 - After: invalidate CPU cache

X-Mas Assignment 13



An unknown MPFC (?) device appeared



☐ Ignore

☒ Write Driver

Task

- You will write a device driver for a MPFC device
- Device Specs
 - Uses single descriptor ring
 - Uses DMA transfers
 - Uses interrupts for “new item” or “no buffer” signals
 - Memory Mapped Device

Steps to do

- Initialize the device (reset) and setting up the in memory data structures
- Start the device
- Start issuing DMA requests to the device
- Activate interrupt and go to sleep
- Do not terminate! Hand back the buffers to the device once used

Hints: Interrupts

- You will need to register a handler which is called for received interrupts
- Interrupt received:
 - Wake the sleeping thread
 - Acknowledge the interrupt

Hints: Buffer Ring

- Allocate enough memory
- Keep track of who owns the buffer
- Tell the device where to find the buffer rings!
 - (Normally you would have to give the physical address, but we stay virtual this time)



- Whishing you a merry Christmas and all the best in the new year!